

Lab 1: Fun with System Calls Report

Part a:

Part a is to implement the existing exit system call to accept and store the exit status of the process that is terminated.

There's two possible methods to do this:

- 1) Change the existing system call signature to void exit(int status)
- 2) Create a new exit system call.

We create a new exit system call exitStat(int) and maintain the exit() function.

Part b:

Part b is to update the existing wait system call to int wait(int*status).

To do this , we created a wait system called int wait(*) to prevent the current process until its child processes successfully exit. The int wait(*) returns the terminated child's exit status and it returns 0 or -1 based on its execution.

Part c:

Part c is to add a waitpid system call int waitpid(int pid, int *status, int options) based on wait system call from part 2. The goal is to implement it so that it waits for a process that has match pid and returns a pid of the terminated process. Waitpid function acts like wait sys call and waits for a process (not necessary a child process) with a pid that equals to the pid provided in the parameter.

The following screenshots below are our modification of the code we made in lab 1

user.h:

```

int open(const char*, int);
struct rtcdate;

// system calls
int fork(void);
int exit(void) __attribute__((noreturn));
int wait(int*); // Updated wait syscall signature to int wait(int *status) - assignment 1
int pipe(int*);
int write(int, const void*, int);
int read(int, void*, int);
int close(int);
int kill(int);
int exec(char*, char**);
int open(const char*, int);
int mknod(const char*, short, short);
int unlink(const char*);
int fstat(int fd, struct stat*);
int link(const char*, const char*);
int mkdir(const char*);
int chdir(const char*);
int dup(int);
int getpid(void);
char* sbrk(int);
int sleep(int);
int uptime(void);
// Added two syscalls below - assignment 1
int exitStat(int);
int waitpid(int, int*, int);

```

defs.h:

```

//PAGEBREAK: 16
// proc.c
int      cpuid(void);
void     exit(void);
int      fork(void);
int      growproc(int);
int      kill(int);
struct cpu* mycpu(void);
struct proc* myproc();
void     pinit(void);
void     procdump(void);
void     scheduler(void) __attribute__((noreturn));
void     sched(void);
void     setproc(struct proc*);
void     sleep(void*, struct spinlock*);
void     userinit(void);
int      wait(int*); // Update the wait system call signature to int wait(int *status) - assignment 1
void     wakeup(void*);
void     yield(void);

int      exitStat(int); // Change the exit system call signature to void exit(int status) - assignment 1
int      waitpid(int, int*, int); // Add a waitpid system call: int waitpid(int pid, int *status, int options) - assignment 1

// swtch.S

```

sysproc.c:

```
// wait syscall - assignment1
int
sys_wait(void)
{
    int *status;
    argptr(0, (void*)&status, sizeof(status));
    return wait(status);
}
```

```
// Update the wait system call signature to int wait(int *status) - assignment 1
// A handler for our new created exit().
// Reads exit status from the user in the command line argument.
// Then calls new created exit() and takes that argument as its parameter.
int
sys_exitStat(void)
{
    int exit_Status;
    if(argint(0, &exit_Status) < 0){
        return -1;
    }
    return exitStat(exit_Status);
}

// A waitpid system call: int waitpid(int pid, int *status, int options) - assignment1
// The system calls a wait for a process (not necessary a child process) with a pid that equals to one provided by the pid argument.
// The return value is the process id of the process that was terminated or -1
// If this process does not exist or an unexpected error occurred.
int
sys_waitpid(void)
{
    int pid;
    int options = 0; // default value
    int* status;
    if(argint(0, &pid) < 0){
        return -1;
    }
    if(argptr(1, (void*)&status, sizeof(status)) < 0){
        return -1;
    }
    return waitpid(pid, status, options);
}
```

proc.h:

```
// Per-process state
struct proc {
    uint sz;                // Size of process memory (bytes)
    pde_t* pgdir;           // Page table
    char *kstack;           // Bottom of kernel stack for this process
    enum procstate state;   // Process state
    int pid;                // Process ID
    struct proc *parent;    // Parent process
    struct trapframe *tf;   // Trap frame for current syscall
    struct context *context; // swtch() here to run process
    void *chan;             // If non-zero, sleeping on chan
    int killed;             // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;      // Current directory
    char name[16];          // Process name (debugging)

    int exitstatus;         // Exit Status - assignment 1
};
```

proc.c:

```
/*
 * wait for a child process to exit and return its pid.
 * Return -1 if this process has no children.
 * int wait(int* status) reads the child's exit status (exitstatus) we defined in struct proc in proc.h
 * The child exit status (exitstatus) is call in proc.c's function exitStat()
 * - assignment 1
 */
int
wait(int* status)
{
    struct proc *p;
    int havekids, pid;
    struct proc *curproc = myproc();

    acquire(&ptable.lock);
    for(;;){
        // Scan through table looking for exit children.
        havekids = 0;
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->parent != curproc)
                continue;
            havekids = 1;
            if(p->state == ZOMBIE){
                // Found one.
                pid = p->pid;
                kfree(p->kstack);
                p->kstack = 0;
                freevm(p->pgdir);
                p->pid = 0;
                p->parent = 0;
                p->name[0] = 0;
                p->killed = 0;
                p->state = UNUSED;
                if(status){
                    *status = p->exitstatus;
                }
            }
        }
    }
    return pid;
}
```

```

        *status = p->exitstatus;
    }
    p->exitstatus = 0;
    release(&ptable.lock);
    return pid;
}
// No point in waiting if we don't have any children.
if(!havekids || curproc->killed){
    release(&ptable.lock);
    return -1;
}
// Wait for children to exit
sleep(curproc, &ptable.lock);
}
/PAGEBREAK: 42
/ Per-CPU process scheduler.

```

316,5 45%

```

/* An exited process remains in the zombie state
until its parent calls wait() to find out it exited.
New exit function with int.
An exit process remains in a zombie state until its parent calls wait().
assignment 1
*/
int
exitStat(int status)
{
    // cprintf("Exit status: %d\n", status);

    struct proc *curproc = myproc();
    struct proc *p;
    int fd;

    if(curproc == initproc)
        panic("init exiting");

    // Close all open files.
    for(fd = 0; fd < NOFILE; fd++){
        if(curproc->ofile[fd]){
            fileclose(curproc->ofile[fd]);
            curproc->ofile[fd] = 0;
        }
    }

    begin_op();
    iput(curproc->cwd);
    end_op();
    curproc->cwd = 0;

    acquire(&ptable.lock);

    // Parent might be sleeping in wait().
    wakeup1(curproc->parent);

    // Pass abandoned children to init.

```

541,5 81%

```

// Pass abandoned children to init.
for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
    if(p->parent == curproc){
        p->parent = initproc;
        if(p->state == ZOMBIE)
            wakeup1(initproc);
    }
}

// Jump into the scheduler, never to return.
curproc->state = ZOMBIE;
curproc->exitstatus = status;
sched();
panic("zombie exit");
}
/PAGEBREAK: 36
/ Print a process listing to console. For debugging.

```

553,5 84%

This method waits for a process (not necessary a child process) with a pid that equals to the one provided by the pid argument.
The return value is the process id of the process that was terminated or -1.
This process does not exist or if an unexpected error occurred.
This function is similar to the wait() defined above.
assignment 1

```
//
int
waitpid(int pid, int* status, int options)
{
    struct proc *p, *curproc = myproc();
    int found_process; //similar to havekids in wait()
    acquire(&ptable.lock);

    //loops continuously till the process with given pid is terminated
    for(;;) {
        found_process = 0;

        //Scan through the process table looking for exited processes. Terminated
        //processes will be in ZOMBIE state.
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
            // If the process pid does not match the given pid, no need to continue
            // with this process.
            if(p->pid != pid) continue;

            found_process = 1;
            if(p->state == ZOMBIE) {
                //Found the process with the given pid that has exited.
                kfree(p->kstack);
                p->kstack = 0;
                freevm(p->pgdir);
                p->pid = 0;
                p->parent = 0;
                p->name[0] = 0;
                p->killed = 0;
                p->state = UNUSED;
            }
        }

        if(found_process) {
            if(status) *status = p->exitstatus;
            p->exitstatus = 0;
            release(&ptable.lock);
            return pid;
        } else if(options == 1) { //if options is passed by the user.

            //the process with the given pid is still running, so we
            //don't block the current process, just release the lock on
            //ptable and return 0.
            release(&ptable.lock);
            return 0;
        }
    }

    // No point waiting if the the process with given pid does not exist
    // or the current process is killed.
    if(!found_process || curproc->killed) {
        release(&ptable.lock);
        return -1;
    }

    // Wait for the process with the given pid to exit.
    sleep(curproc, &ptable.lock);
}
```

629,1-8 95%

```

        p->killed = 0;
        p->state = UNUSED;
        if(status) *status = p->exitstatus;
        p->exitstatus = 0;
        release(&ptable.lock);
        return pid;
    } else if(options == 1) { //if options is passed by the user.

        //the process with the given pid is still running, so we
        //don't block the current process, just release the lock on
        //ptable and return 0.
        release(&ptable.lock);
        return 0;
    }
}

// No point waiting if the the process with given pid does not exist
// or the current process is killed.
if(!found_process || curproc->killed) {
    release(&ptable.lock);
    return -1;
}

// Wait for the process with the given pid to exit.
sleep(curproc, &ptable.lock);
}
```

654,1 Bot

.c files needed to modified:

Many .c files required a change from exit() to exit(0) and wait() to wait(0).

The following files were changes that we made:

cat.c echo.c forktest.c grep.c init.c kill.c ln.c ls.c mkdir.c rm.c

sh.c stressfs.c trap.c ulib.c wc.c zombie.c usertest.c

Test:

To run our test file; lab1.c

Use make clean qemu-nox inside the xv6 directory.

Then run: Lab1 1 to test part a and b.

Finally, Lab1 2 to test part c.

However, a weird format issue occurred for the part c output as shown below.

```
vagrant@ubuntu-xenial:~/lab/CS153/Lab1$ make clean qemu-nox
```

```
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ lab1 1
```

```
This program tests the correctness of your lab#1
```

```
Parts a & b) testing exit(int status) and wait(int* status):
```

```
This is child with PID# 4 and I will exit with status 0
```

```
This is the parent: child with PID# 4 has exited with status 0
```

```
This is child with PID# 5 and I will exit with status -1
```

```
This is the parent: child with PID# 5 has exited with status -1
```



```

$ lab1 2

This program tests the correctness of your lab#1

Part c) testing waitpid(int pid, int* status, int options):

This is the child with PID# 10 and I will exit with status 11
This is the child with PID# 8 and I will exit with status 12
This is the child with PID# 9 and I will exit with status 13
This is the child with PID# 14
Parent: Now waiting for child with PID# 10

This is the parent: Now waiting for child with PID# 8
Parent: Child# 8 has exited with status 12

This is the parent: Now waiting for child with PID# 9
Parent: Child# 9 has exited with status 13

This is the parent: Now waiting for child with PID# 7
Parent: Child# 7 has exited with status 11

This is the parent: Now waiting for child with PID# 11
Parent: Child# 11 has exited with status 15

```

Below, we have the diff output txt file.

diff output txt file: