

Lab 2: Scheduler

Implement Priority Scheduling:

The purpose of lab 2 is to implement a scheduler from a simple round robin to a priority scheduler. We added a new system call that changes the priority of a process in which process with the highest priority (the one with the lowest value) gets scheduled first. The system call we added: `int setPriority(int p)` and `void getPriority(void)`.

- 1) `int setPriority (int priority)`: The current process is obtained by using `myproc()`. The process priority is updated with the passing argument.
- 2) `int getPriority()`: To print updated priorities , we created a getter system call

We also have to modified the following files: `syscall.h`, `syscall.c`, `defs.h`, `user.h`, `sysproc.c`, `usys.S`, `proc.h` and `proc.c`

Adjust priority (avoid starvation):

In order to avoid starvation , we implement our aging of priority by decreasing priority of processes due to high priority and increased priority of processes which to run.

We do this by modifying the `proc.c` scheduler function; `scheduler(void)`. The change in the scheduler function from simple round robin to a priority scheduler.

Tracking Scheduling Performance:

We created a turnaround time for each process to print when the process exists. To track the scheduling performance of each process, we modified the `proc.c` exit function; `exit(void)`. The change is needed to print the turnaround time when it exists.

The changes in our code are listed below.

syscall.h:

```
#define SYS_setPriority 24
#define SYS_getPriority 25
```

syscall.c:

```
// Added syscall to set and get priority of a process - assignment 2
extern int sys_setPriority(void);
extern int sys_getPriority(void);
```

```
[SYS_setPriority]    sys_setPriority,
[SYS_getPriority]    sys_getPriority,
};
```

usys.S:

```
SYSCALL(setPriority)
SYSCALL(getPriority)
```

defs.h:

```
int    setPriority(int); // Add set Priority function - assignment 2
int    getPriority(void);
```

sysproc.c:

```

int sys_setPriority(void)
{
    int priority;
    if (argint(0,&priority)<0){
        return -1;
    }
    return setPriority(priority);
}

int sys_getPriority(void)
{
    return getPriority();
}

```

user.h:

```

// Add user syscall for setPriority - assignment 2
int setPriority(int);
int getPriority(void);

```

proc.h:

```

// Per-process state
struct proc {
    uint sz;                // Size of process memory (bytes)
    pde_t* pgdir;           // Page table
    char *kstack;           // Bottom of kernel stack for this process
    enum procstate state;   // Process state
    int pid;                // Process ID
    struct proc *parent;    // Parent process
    struct trapframe *tf;   // Trap frame for current syscall
    struct context *context; // swtch() here to run process
    void *chan;             // If non-zero, sleeping on chan
    int killed;             // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;      // Current directory
    char name[16];          // Process name (debugging)
    int exitstatus;         // Exit Status - assignment 1
    int priority;           // Priority value - assignment 2
    int start_time;
};

```

proc.c

```

void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;
    int low_priority;

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        low_priority = 1000;
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if (p->state == RUNNABLE && p->priority < low_priority) {
                low_priority = p->priority;
            }
        }

        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

            if(p->priority != low_priority){
                if (p->priority > 0) {
                    p->priority--;
                }
                continue;
            }
        }

        // Switch to chosen process.  It is the process's job
        // to release ptable.lock and then reacquire it
        // before jumping back to us.
        c->proc = p;
        switchvm(p);
        p->state = RUNNING;
        p->priority++;
    }
}

```

```

        p->state = RUNNING;
        p->priority++;

        swtch(&(c->scheduler), p->context);
        switchkvm();

        // Process is done running for now.
        // It should have changed its p->state before coming back.
        c->proc = 0;
    }

    release(&ptable.lock);
}
}

```

```
//get and set priority of process - assignment
int setPriority(int priority)
{
    struct proc *p = myproc();
    p->priority = priority;
    return 0;
}

int
getPriority()
{
    struct proc *curproc = myproc();
    return curproc->priority;
}
```

```
curproc->cwd = 0;

end_time = ticks;
cprintf("\n turnaround time is %d\n", end_time - curproc->start_time);
```

Lab 2 Test Code:

```

#include "types.h"
#include "user.h"

int main(int argc, char *argv[])
{
    int PScheduler(void);
    printf(1, "\n This program tests the correctness of your lab#2\n");
    PScheduler();
    return 0;
}

int PScheduler(void){
    // use this part to test the priority scheduler. Assuming that the priorities range between 0 to 31
    // 0 is the highest priority and 31 is the lowest priority.
    int pid;
    int i,j,k;
    int priorityArr[] = {30, 15, 25, 0, 0};

    printf(1, "\n Testing the priority scheduler and setPriority(int priority) system call:\n");
    printf(1, "\n Assuming that the priorities range between 0 to 31\n");
    printf(1, "\n 0 is the highest priority. All processes have a default priority of 10\n");
    printf(1, "\n The parent processes will switch to priority 0\n");
    setPriority(0);

    for(i = 0; i < 5; i++) {
        pid = fork();
        if (pid > 0) {
            continue;
        } else if (pid == 0) {
            setPriority(priorityArr[i]);
            printf(1, "\n child# %d has priority %d before starting its work", getpid(), getPriority());
            for (j=0;j<50000;j++) {
                for(k=0;k<1000;k++) {
                    asm("nop");
                }
            }
            printf(1, "\n child# %d has priority %d after finishing its work", getpid(), getPriority());
            printf(1, "\n child# %d with original priority %d has finished! \n", getpid(), priorityArr[i]);
        }
    }
}

```

```

        exit();
    } else {
        printf(2, "\n Error \n");
    }
}

if(pid > 0) {
    for(i = 0; i < 5; i++) {
        wait(0);
    }
    printf(1, "\n if processes with highest priority finished first then its correct \n");
}
exit();
}

```

To run our test file; lab2.c

Use make clean qemu-nox inside the xv6 directory.

Enter Lab2 to run our test code program.

However, a weird format issue occurred for the test output as shown below.

```
vagrant@ubuntu-xenial:~/CS153/Lab2$ make clean qemu-nox
```



```
173524 bytes (174 kB, 169 KiB) copied, 0.00142049 s, 122 MB/s
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw -drive file=xv6.img,index=0,media=disk,format=raw
xv6...
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ lab2
```

Test output:

```
$ lab2

This program tests the correctness of your lab#2

Testing the priority scheduler and setPriority(int priority) system call:

Assuming that the priorities range between range between 0 to 31

0 is the highest priority. All processes have a default priority of 10

The parent processes will switch to priority 0

child# 4 has priority 3
child# 5 has priority 15 before starting its work
child# 6 has priority 25 before starting its work
child# 7 has priority 0 before starting its work
child# 8 has priority 9 before starting its work
child# 7 has priority 7 after finishing its work
child# 7 with original priority 0 has finished!

turnaround time is 19
0 before starting its work
child# 8 has priority 6 after finishing its work
child#
child# 5 has priority 6 after finishing its8 with original priority 9 has finished!

turnaround time is 43
work
child# 5 with original priority 15 has finished!
```



```
turnaround time is 43
work
child# 5 with original priority 15 has finished!

turnaround time is 52

child# 6 has priority 5 after finishing its work
child# 6 with original priority 25 has finished!

turnaround time is 58

child# 4 has priority 9 after finishing its work
child# 4 with original priority 30 has finished!

turnaround time is 62

if processes with highest priority finished first then its correct

turnaround time is 72
*
```